

DYNAMIC FINGER AND WORKING SET: APPROACHING THE LOWER BOUNDS

Luis Barba, Rolf Fagerberg and Pat Morin

August 22, 2013

ABSTRACT. This paper begins the process of making practical dynamic comparison-based dictionaries that have the dynamic finger and working-set properties. In particular, we initiate the study of the exact number of comparisons needed in dynamic dictionaries that have these two properties. We present a data structure with the dynamic finger property that performs at most $\log_2 d + o(\log d)$ comparisons per operation and a data structure that, for any $\varepsilon > 0$, has the working-set property and performs at most $(1 + \varepsilon)\log_2 w + o(\log w)$ comparisons per search. Here, d and w are the “finger number” and “working-set number”, respectively, of the element being accessed. All previously-known structures with the dynamic-finger or working-set property perform at least $2\log_2 d$ or $4\log_2 w$ comparison, respectively.

1 Introduction

In this paper we present comparison-based dictionaries that support finger-search and that have the working-set property and use a nearly-optimal number of comparisons per search. In the remainder of this introduction we define these terms and explain why these results are interesting.

1.1 Comparison-Based Dictionaries

Comparison-based dictionaries supporting the three *basic operations* insert, delete and search represent the classic data-structuring problem in computer science. AVL trees, which support all three basic operations with an asymptotically optimal running time were discovered already in 1962 [?]. AVL trees implement the basic operations in $O(\log n)$ time per operation while performing at most $1.4404 \log(n + 2)$ comparisons between the element being inserted/deleted/searched and the elements stored in the AVL tree. (Here, and throughout, $\log x$ is a shorthand for $\log_2 \max\{2, x\}$.)

Since the discovery of AVL trees, many other comparison-based dictionaries have been proposed that perform at most $c \log n + O(1)$ comparison per operation, for some constant $c > 1$: 2-3 trees ($c = XX$) [?], red-black trees ($c = 2$) [?], splay trees ($c = 3$) [?], scapegoat trees ($c = 1 + \epsilon$) [?, ?], skiplists ($c = e/\log e \approx XX$) [?, ?], treaps ($c = 2/\log e \approx XX$) [?, ?], and randomized binary search trees ($c = 2/\log e \approx XX$) [?] are just a few of the colorful names for data structures that support the basic operations in $O(\log n)$ time per operation.¹

A handful of lesser-known structures reduce the constant c to 1 while still supporting all operations in $O(\log n)$ time: Lai, Andersson, and Fagerberg [?, ?, ?] each present data structures supporting each of the basic operations in $O(\log n)$ (amortized) time using at most $\log n + O(1)$ comparisons. Since $\log(n + 1)$ is an information-theoretic lower-bound on the *expected* number of comparisons when searching for a random key, these results more or less close book on the worst-case number of comparisons achievable by comparison-based dictionaries that support all operations in $O(\log n)$ time.

1.2 Dictionaries with Properties

A more recent trend in the design of comparison-based dictionaries is the study of dictionaries with special running-time *properties*. These results are motivated by a practical observation: in real applications, queries are usually not independently and uniformly distributed, so the information-theoretic lower bound does not apply. Data structures having properties that exploit special patterns in query sequences do better than $\Theta(\log n)$ time per query. Examples of such properties include:

the static-optimality property: in which the average access time is proportional to the (empirical) entropy of the distribution of searches.

The study of *biased dictionaries*, that satisfy the static-optimality property has a long and rich history. It is known that, given a set of keys and a distribution of queries,

¹Some of these structures use amortization, in which case the $O(\log n)$ running-time bound is only an amortized bound and some use randomization, in which case the running-time bound is an expected running time bound, with the expectation being taken over random choices made within the data structure.

it is possible to construct, in quadratic time, a binary search tree that minimizes the expected number of comparison per search [?]. A nearly optimal binary search tree, that does at most two extra comparisons per search can even be constructed in linear time [?].

the dynamic finger property: in which the time to perform the current search for an element, x , is logarithmic in the difference, d , in ranks between the element x being searched for and the element, x_0 , returned as a result of the previous search. (During the first access, d is defined to be n .)

Several data structures are known that can perform searches in $O(\log d)$ time using $c \log d + o(\log d)$ comparisons; these include homogeneous finger search trees ($c = 4$) [?], splay trees ($c = 250,000$) [?], treaps ($c = 4/\log e$) [?], and skiplists ($c = 2e/\log e$) [?], and unary-binary trees augmented with hands ($c = 2$) [?, ?]. See Brodal [?] for a recent survey on finger search.

the working-set property: in which the time to perform the current search for an element, x , is logarithmic in the number, w , of distinct elements accessed since the most recent previous access to x . (If this is the first access to x , then w is defined to be n .)

Several data structures are known that can perform searches in $O(\log w)$ time using $c \log w + o(\log w)$ comparisons; these include splay trees ($c = 4$) [?],² several variants of Iacono's doubly-exponential structure ($c \geq 4$) [?], variants of skiplists and B-trees ($c = ??$), and layered working-set trees ($c = ??$) [?].

It is worth noting that the working-set property is stronger than the static optimality property. Roughly speaking, any data structure that has the working-set property with constant c can perform a sufficiently long sequence of searches with an average of $c\tilde{H}$ comparisons per search, where \tilde{H} is the empirical entropy of the search sequence [?, ?]. Note that, unlike the results for static optimality, this does not require advanced knowledge of the query distribution. This observation, which follows from Jensen's Inequality, is the basis of one of the most effective and practical data compression algorithms known [?, ?].

Additional properties have been proposed, including the unified property [?] and the dynamic optimality property [?], but the three properties discussed above are among the oldest seem to be among the most fundamental.

1.3 Constants for Dynamic-Finger and Working-Set

Unfortunately, dictionaries with properties generally fail to deliver on the practical results they promise, and this is due to the constants in their running times. To illustrate this, consider a binary search tree, T , that contains $n = 10^6$ elements, and is implemented using Andersson, Lai, or Fagerberg's binary search trees. Such a structure can perform any search using at most

$$\lceil \log(10^6 + 1) \rceil + 1 = 21$$

²Tarjan and Sleator's Working Set Lemma [?] gives the bound $c = 8$. However changing their potential from $1/w(i)^2$ to $1/w(i) \log^2 w(i)$ improves this to $c = 4$.

comparisons. In contrast, consider storing the same elements in a dictionary, S , implemented using splay trees. The (amortized) number of comparisons performed during a search in S is $4\log w + o(\log w)$ (see Section ??). Thus, in order for S (a splay tree) to perform faster than T (a balanced binary search tree), we must have

$$4\log w < 21 \Leftrightarrow \log w < 5.25 \Leftrightarrow w \leq 38 .$$

That is, of the one million elements stored in S and T , there are only 38 that can be accessed faster in S than in T . Most of the remaining $10^6 - 38$ elements require more comparison to access in S than in T , and the vast majority require a factor of 2–4 times as many comparisons.

To further widen the performance gap, the splay tree, S , performs $3\log w$ rotations after every search, while searches do not affect the structure of T at all. Taking all this into consideration, it seems difficult to imagine *any* realistic application where the working-set property of splay trees would give better performance than (Andersson, Lai, or Fagerberg’s) balanced binary search trees. (This back-of-the-envelope analysis is born out by several experimental evaluations of splay trees [?, ?, ?].)

1.4 New Results

So it seems that a necessary condition for a dynamic-finger structure or a working-set structure to be practical is that the leading constant on the number of comparisons performed during a operation should be very small; ideally this constant should be 1. In this paper we present two data structures that attempt to achieve this ideal. Our first data structure supports dynamic-finger operations in $O(\log d)$ time and performs at most $\log d + o(\log d)$ comparisons during any operation. Our second data structure is parameterized by a parameter $\varepsilon > 0$ and supports all operations in $O(\varepsilon^{-1} \log w)$ time and performs at most $(1 + \varepsilon)\log w + o(\log w)$ comparisons during any operation. (Note that ε may even depend on n .)

1.5 Focusing on Comparisons

We do not claim that minimizing comparisons immediately implies that our data structures are efficient in practice. Like all dictionaries, our algorithms perform operations other than comparisons, so justifying such a claim would require some algorithms engineering to streamline and tune our data structures as well as a rigorous experimental evaluation. Our only claim is that our data structures satisfy a necessary condition that prevents existing structures, such as splay trees, from being practical.

Nevertheless there are programming environments that tilt the table heavily towards the goal of minimizing comparisons. Consider, for example, the Java Collections Framework [?]. In the JCF, a single comparison between two Integer objects stored in a SortedMap (which is implemented as a red-black tree) involves

1. dereferencing the data structure’s Comparator object, c (`this.c`);
2. calling c ’s `compare(a,b)` method, which requires both a function call and a dereferencing operation in c ’s dispatch table (`c.compare(a,b)`); and

-
3. the `compare(a,b)` method must then dereference `a` and `b`'s native `int` variables and compare them by computing their difference using the `isub` instruction. (`return a.value - b.value`)³

In this way, a single comparison performed within the data structure involves a function call and four dereferencing operations, each of which is an order of magnitude slower than the `isub` instruction that actually performs the comparison. One could argue that this is a(nother) reason not to develop performance-critical software in Java, but given that there are already 3 billion devices running Java [?], it is still a worthwhile goal to optimize algorithms for it.

2 The Dynamic-Finger Property

Other than splay trees (which have an enormous constant), data structures that perform finger search typically work in two phases.⁴ In the *upward phase* the search starts at the current node, which contains x_0 , and walks upwards in the data structure until reaching a node, w , from which the usual search procedure can be applied. In the *downward phase* a search for x that starts at w is performed.

This two phase approach is expensive: If the worst-case number of comparisons done during a (non-finger) search is $c \log n + o(\log n)$ the two phase approach usually leads to a search time of $2c \log d + o(\log d)$; using a simple two-phase finger search on a normal structure doubles the leading constant. Informally, this is because the logarithm is such a slow-growing function: If w is a node containing a value equidistant from x_0 and x , then the the number of comparison used to go from x_0 to w and then to x is $c \log(d/2) + c \log(d/2) = 2c \log d - 2c$. This is the reason no existing structure has a constant smaller than 2.

We note that there are several approaches to resolve this problem. One particularly easy method is to modify the upward phase of a two phase algorithm so that it only performs $O(\log \log d)$ comparisons. This modification involves a straightforward implementation of exponential search (discussed below). In typical existing structures, all of which are pointer based, this modification will decrease the number of comparisons during the upward phase, but the running time will still be $O(\log d)$. (This is analogous to implementing binary search on a linked list; the number of comparisons is logarithmic, but the running time is still linear.)

2.1 Overview

We will describe two phase approach to finger search in which the first phase runs in $O(\log \log d)$ time and performs $O(\log \log d)$ comparisons and the second phase runs in $O(\log d)$ time and performs $\log d + O(1)$ comparisons.

At a high level, our new structure works as follows: We use an extremely well-

³In Java, the `compare(a,b)` method should return a negative value if $a < b$, zero if $a = b$ and a positive value if $a > b$.

⁴Even splay trees can be viewed as having a two-phase approach to finger search; the first phase (splaying) takes place while searching for the previous element).

balanced search tree, T , which has height $\log n + O(1)$. We then use an auxiliary structure of size $O(\log n)$ that keeps track of information about the search paths around x_0 . This structure is an array-based variant of Blleloch et al.'s *hands* data structure. Our version of *hands* supports two phase finger searches with the first phase running in $O(\log \log d)$ time and the second phase running in $O(\log d)$ time.

Of course, there is an issue with this approach that needs to be resolved. Blleloch et al.'s *hands* are a pointer-based structure that make use of the fact that linked lists can be split and joined in constant time. To obtain our speedup of the first phase, we require an array-based structure. This is a problem since, in general, arrays can not be split and joined in constant-time. We show that this problem can be handled effectively using an indexing trick. As a side-effect, we obtain an array-based version of Blleloch et al.'s *hands* that is (arguably) easier to maintain and likely to be more efficient than the original version.

2.2 Tools

We begin with a short discussion of the three tools used by our algorithm.

2.2.1 Unbounded Search in a Sorted Array

A standard method for searching in a sorted array, $A = a_1, \dots, a_n$, when there is reason to believe that the search result is closer to the front of the array than the back is *exponential search*. To find the position of some query value, x , we examine a_1, a_2, a_4, a_8 , and so on until finding the first value, $j \geq 0$, such that $a_{2^j} \geq x$. At this point, a standard binary search over $a_{2^{j-1}+1}, \dots, a_{2^j}$ can be used to find x using an additional j comparisons.

Exponential search uses $2j$ comparisons. On the other hand, the index, i , of x , is greater than $2^{j-1} + 1$, so the number of comparisons performed by exponential search, when the search result is at position i , is at most $2\lceil \log i \rceil \leq 2\log i + 1$. The idea behind exponential search has been refined by Bentley and Yao [?]. Using their algorithm one can, for example, find x using $\log i + O(\log \log i)$ comparisons.

2.2.2 1-2 Trees

To obtain a dynamic dictionary, we will use a special kind of very well balanced binary search trees. A *unary-binary tree* or *1-2 tree* is a tree in which all the leaves have the same depth and every internal node has one (a unary node) or two (a binary node) children. The following result, due to Fagerberg [?], is used in the maintenance of binary trees of very low height:

Theorem 1 (Fagerberg). *It is possible to maintain a unary-binary search tree, T , under the operations of insertion and removal in constant amortized time per operation so that, after every operation*

1. (low height) T has height $\log n + O(1)$, where n is the number of leaves currently in the tree;
2. (exponential size) every node of height h is the root of a subtree having at least $c2^h$ leaves, for some constant $c > 0$; and




Figure 1: Btleloch et al.’s hands: The search path for x_0 is in orange. Every node is contained in L and/or R . The stacks L^+ and R^+ contain the large green nodes.

3. (no adjacent unary) the child of each unary node is either a leaf or a binary node.

2.2.3 Hands

Btleloch et al. introduce a data structure, that they call “hands” that augments a degree-balanced search so that it supports efficient finger search. Hands consist of four lists. The first two of these lists are called L and R . Suppose the most recently accessed value was x_0 . To understand the contents of L and R , consider the standard drawing of a binary search tree, where the x coordinate of each node is given by its key and the y -coordinate is given by its height. If we draw a vertical line through x_0 , then L contains those nodes immediately to the left of (and on) this line and R contains those nodes immediately to the right of (and on) this line (see Figure 1). (Informally, L and R contain the search paths for $x_0 - \epsilon$ and $x_0 + \epsilon$.) The nodes on these paths are ordered from bottom-to-top, so that $L[i]$ and $R[i]$ are nodes of height i .

In addition to the two lists L and R , two stacks, L^+ and R^+ , also implemented as lists, are maintained. L^+ contains those nodes on the search path for x_0 that are less than or equal to x_0 and R^+ contains those nodes that are greater than or equal to x_0 . These stacks are ordered so that the nodes closer to x_0 are closer to the top of the stack and x_0 itself is the top element on each the stack. Note that the elements of L^+ and R^+ are a subset of




Figure 2: Performing a finger search using a hand.

those in L and R , respectively.

Finger Search. Refer to Figure 2. Suppose, without loss of generality, that we want to perform a search for some value $x > x_0$. To do this, we repeatedly examine the top element of R^+ and pop it off if it is less than x . This process pops off at least one element, since $x_0 < x$. Consider the last element \hat{x} popped off the stack. It must be the case that $x = \hat{x}$ or x is in the right subtree of \hat{x} . Unfortunately, searching in \hat{x} 's subtree may be too slow; the height of this subtree has no relation to d .

Consider the second-last element of \hat{x} that was popped off the stack (if no such element exists, imagine \hat{x} is the external node, of height -1 , that is just to the right of x_0). Now, \hat{x} has some height, h , and therefore \hat{x} 's right subtree has size at least $c2^h$. All the elements in \hat{x} 's right subtree are in the interval (x_0, x) , so $d \geq c2^h$. In other words, we can afford to spend $O(h)$ time to find x .

The list R contains \hat{x} as well as a vertex x^* that is one level above \hat{x} . The vertex x^* is either equal to \hat{x} or is in \hat{x} 's right subtree. By starting at x^* and walking upwards until reaching a node whose value is greater than x , we find a node, u , of height $h+r$ that is an ancestor of the node x we are searching for. Since u is the first such node, we also know that $d > c2^{h+r-1}$, so we can afford to search for x starting from node u .

Updating the Hands. In this way, we find x in $O(h+r)$ time using at most $2(h+r)$ comparisons. What remains is to update L , R , L^+ , and R^+ . One can observe that L and R only change at indices $0, \dots, h+r$ and updating $L[0], \dots, L[h+r]$ and $R[0], \dots, R[h+r]$ is easy to do in $O(h+r)$ time starting from the node u .

The tricky part is updating L^+ and R^+ . Updating L^+ is the easier of the two. It is first truncated so that it doesn't contain any nodes at height lower than the height of \hat{x} (Blelloch et al. manage this by maintaining cross pointers between elements of L and R at the same level as well as cross pointers between elements of L and L^+ and R and R^+). The stack L^+ is then extended by adding the appropriate nodes on the search path from u to x .

To update R^+ Blelloch et al. make use of the fact that lists are concatenable. The sequence of nodes that need to be added to R^+ are those on the search path from \hat{x} to x . The portion of this path from \hat{x} to u can be spliced from the sublist $R[h+r], \dots, R[\text{height}(\hat{x})]$. The portion from u to x can be added to R^+ while performing the search from u to x .

2.3 Fast Hands

To speed up Blelloch et al.'s hands, we implement the lists L and R as arrays. The lists L^+ and R^+ are also implemented as arrays, but not in the obvious way. Recall that L^+ and R^+ store nodes on the search path for x_0 , and these nodes are also stored in L and R , respectively.

Observe that the search path for x_0 alternately takes subpaths of L and subpaths of R . Therefore, we can represent L^+ and R^+ as a sequence of pairs of indices, where the pair (ℓ, h) in, for example, R^+ indicates that the stack represented by R^+ contains the elements $R[\ell], \dots, R[h]$. We can even implement R^+ and L^+ as a single array, with odd indices for R^+ and even indices for L^+ ; this combined array then represents the entire search path for x_0 . See Figure ??.

We will use the notation $R_{[]}^+$ and $L_{[]}^+$ to denote the array of pairs that represents the stack R^+ and L^+ , respectively. We observe that this representation of R^+ and L^+ still allows for exponential search. To perform a search on R^+ , for example, we perform exponential search on $R_{[]}^+$ to find an interval (ℓ, h) that contains the element we are searching for. We then perform exponential search, within L and starting at position ℓ , to find the element in L that we are searching for. If the element we are searching for is at distance i from the top of the stack R^+ , then it is straightforward to verify that double application of exponential search takes $O(\log i)$ time and performs $O(\log i)$ comparisons.

Finger Search. Recall that a finger search for $x > x_0$ first searches R^+ in order to find \hat{x} and then searches R in order to find u . In our array-based representation, the first of these searches can be done in $O(\log h)$ time and the second search can be done in $O(\log r)$ time using exponential search. Since $d \geq c2^{h+r}$, this implies that the upward phase of the algorithm takes $O(\log \log d)$.

Updating the Hands. Updating L and R during a finger search is easy. These arrays only need to be updated at positions $0, \dots, h+r$ and this can be done starting u , in $O(h+r)$ time (see Blelloch et al. [?] for details).

To update L^+ , we first need to truncate it so that it does not contain any nodes lower than hand \hat{x} . Note that we have located \hat{x} in $R_{[]}^+$. In particular, we have found an index i such that $R_{[]}^+[i]$ is an interval of R^+ that contains \hat{x} . This means we want to truncate L^+ at location i or $i+1$, so this can be done in constant time. The only other update required on




Figure 3: An array-based implementation of hands.

L^+ is to add nodes on the search path from u to x , which can be done while traversing this path.

To update R^+ we first truncate R^+ at \hat{x} , which is easily done in constant time, since we have already located the interval of R_{\square}^+ that contains \hat{x} . Next, we need to add the path from the right child of \hat{x} to u . This can also be done in constant time by appending the pair $(h+r, \text{height}(\hat{x}) - 1)$ to R_{\square}^+ . Finally, we add the appropriate elements on the search path from u to x while traversing this path.

In summary, it is possible to perform a finger search and update the hands in $O(\log d)$ time using $\log d + O(\log \log d)$ comparisons.

Insertion and Deletion. 1-2 trees support insertion and deletion in constant amortized time provided that a pointer to the node being deleted or a pointer to the location of the insertion is given. The appropriate pointer is easily obtained using finger search.

During an insertion or deletion, the 1-2 tree is modified. Updating the hands during these modification can be done in time proportional to the number of modifications (see Blelloch et al. for details). Therefore, updating the hand does not increase the asymptotical running-time of insertions or deletions. The following theorem summarizes our result for finger search

Theorem 2. *There exists a linear-sized comparison-based dictionary that supports searching in $O(\log d)$ worst-case time using at most $\log d + O(\log \log d)$ comparisons. Insertions and deletions in this data structure involve a single search plus a constant amortized amount of restructuring that can be done in constant amortized time.*

Finally, it is worth noting that our structure maintains all the advantages of hands: It has only $O(\log n)$ size; it does not require any special augmentation (such as level links) of the underlying search tree T , and several processes can share the same search tree, T , each maintaining its own hand. For binary search tree fetishists, this structure can even be implemented as a binary search tree by collapsing unary nodes (see Fagerberg [?] for details).

3 The Working-Set Property

Our working set data structure is a variant of a skiplist. It consists of a sequence of lists L_0, \dots, L_k where

1. $|L_0| \in O(1)$
2. every element in L_i appears in L_{i+1} for all $i \in \{0, \dots, k-1\}$;
3. L_i contains every element, x , such that $w(x) \leq (2 - \epsilon)^i$;
4. if x and y are two consecutive elements of L_i , then at least one of x and y appears in L_{i-1} .

3.1 Searching

A search in a

4 Discussion

4.1 Two-Way Comparisons

4.2 Handling Searches for Missing Values

Acknowledgement

The authors of this paper are partly funded by NSERC and CFI.

References

Authors



Luis Barba. Département D'Informatique, Université Libre de Bruxelles and School of Computer Science, Carleton University

Rolf Fagerberg. Department of Mathematics and Computer Science, University of Southern Denmark

Pat Morin. School of Computer Science, Carleton University